

CS 5/7320

Artificial Intelligence

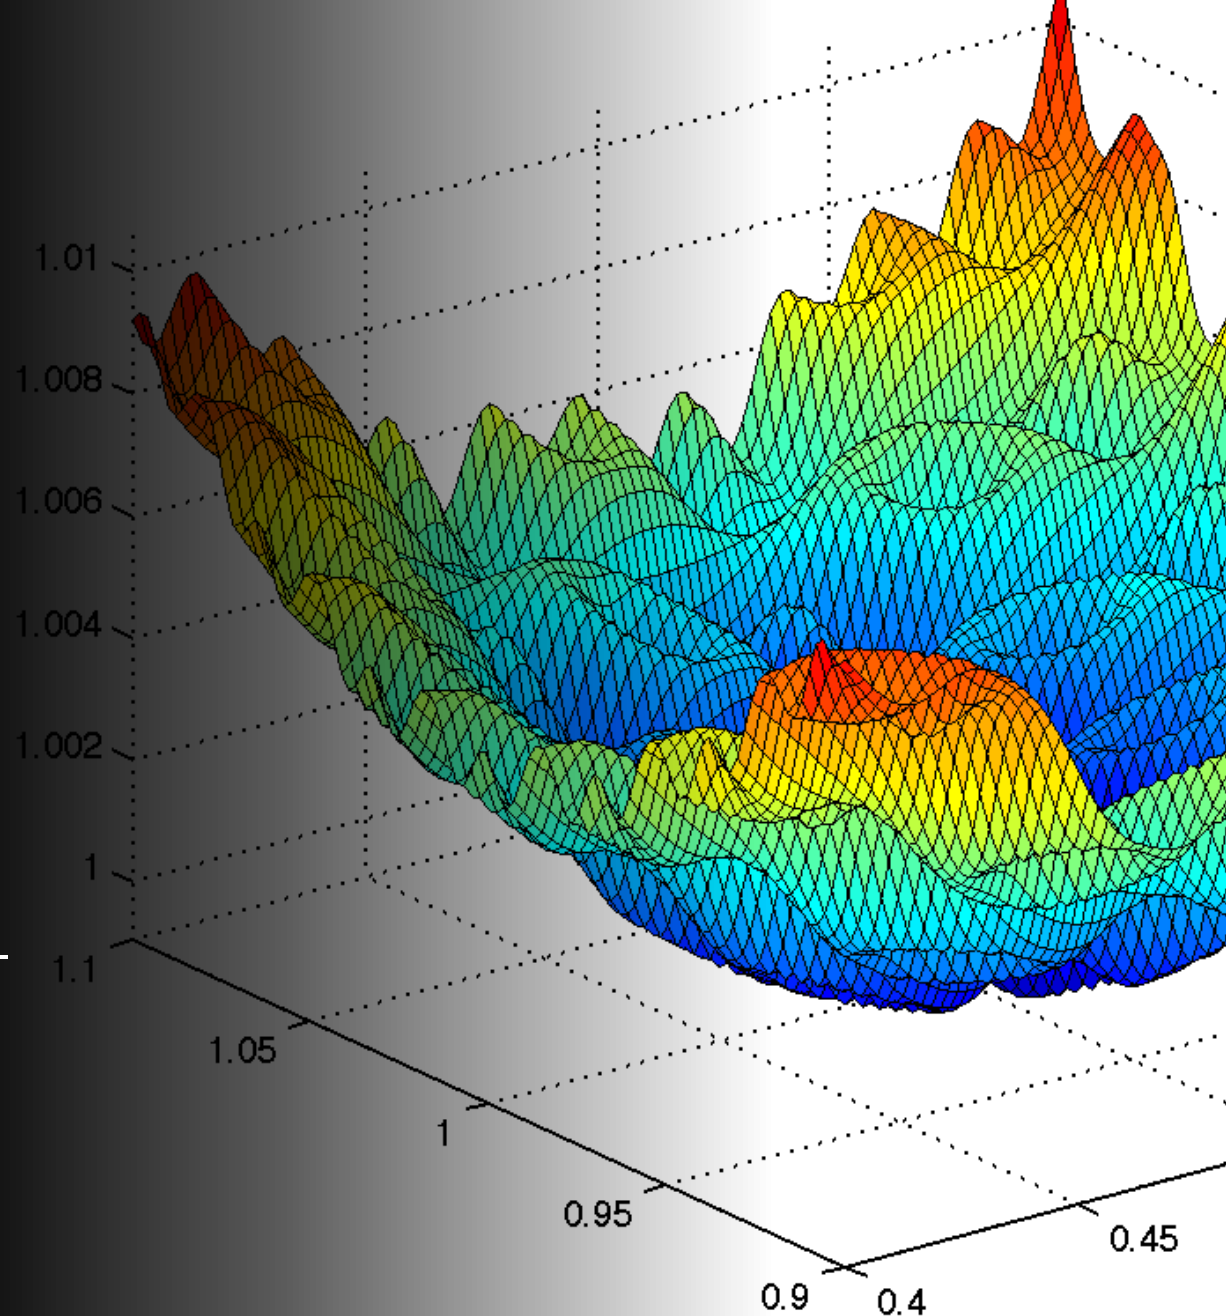
Local Search

AIMA Chapters 4.1 & 4.2

Slides by Michael Hahsler
based on slides by Svetlana Lazepnik
with figures from the AIMA textbook.



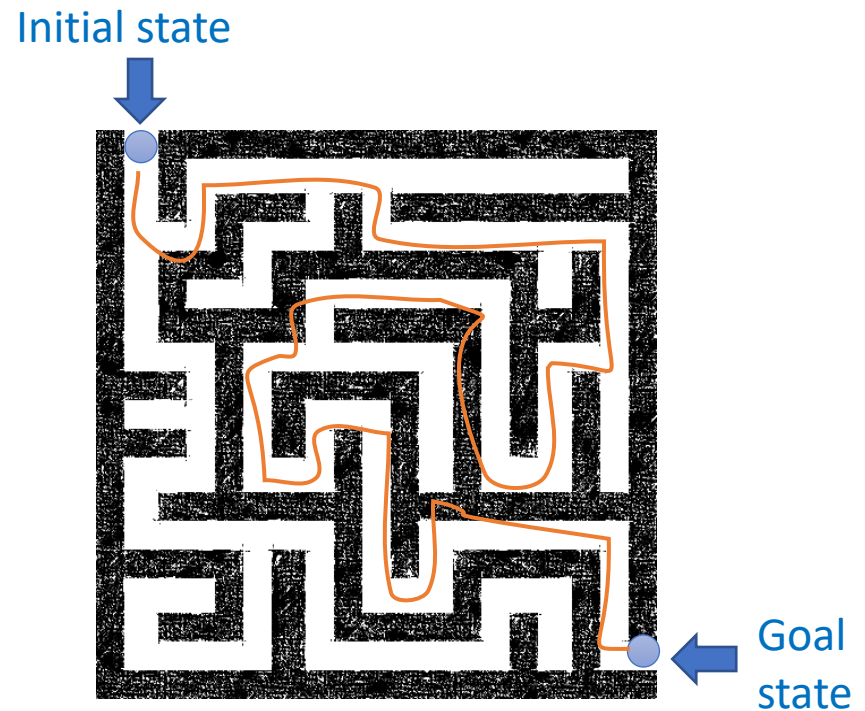
This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).




Recap: Uninformed Search/informed search

Tries to find the
best path
from a
given initial state
to a
given goal state.

- Typically searches a large portion of the search space (needs time and memory).
- Often comes with optimality guarantees.



Local search algorithms



Idea:

- Improve the current solution by moving to a neighboring better state (a.k.a. performing a local move).
- This is fast and needs little memory (no search tree).

- What if we do not know the goal state, but how “good” different states are? This is typically called the **objective function** → finding the best state is an **optimization problem**.
- We need a fast and memory-efficient way to find **the best/a good state**.

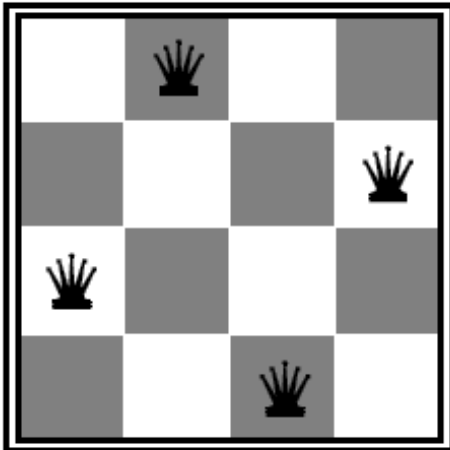
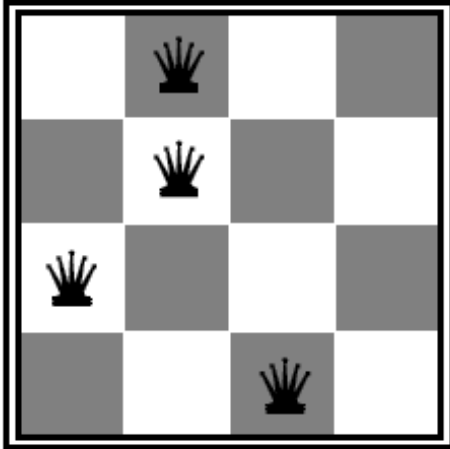
Local search algorithms

Difference to search from the previous chapter:

- a) **Goal state is unknown** and needs to be identified.
- b) Often no explicit initial state + **path to goal and path cost are not important.**
- c) **No search tree.** Just stores the current state.

Use in AI

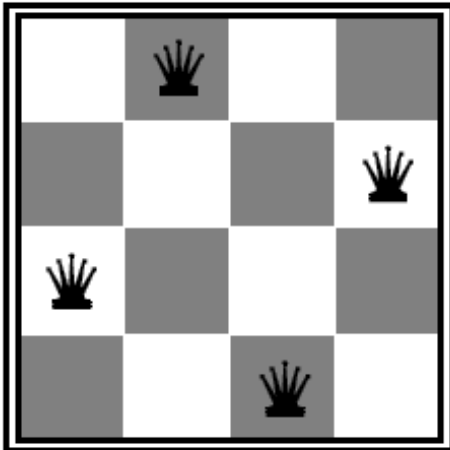

- **Utility-based agent:** Use utility as the objective function and always move to higher utility states. A greedy method used for complicated/large state spaces or online search.
- **Goal-based agent:** Identify a good goal state with a good objective function value before planning the path to that state.
- **General optimization:** Use for effective heuristic search in large or continuous spaces (with an infinite state space). E.g., learn neural networks.



Example: n -queens problem

- **Goal:** Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal.
- **State space:** All possible n -queen configurations. **How many are there?**
- What is a possible **objective function**?

2 conflicts



0 conflicts

Example: n -queens problem

- **Goal:** Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal
- **State space:** all possible n -queen configurations:
4-queens problem: $\binom{16}{4} = 1820$
- What is a possible **objective function**?

Minimize the number of pairwise conflicts


Note: this can be seen as a heuristic used in informed search, but it may not be an admissible heuristic.



Example: Traveling salesman problem

- **Goal:** Find the shortest tour connecting a given set of cities
- **State space:** all possible tours (states are not individual cities!)
- **Objective function:** minimize the length of the tour

Note: We have solved a different problem with uninformed/informed search! Each city was defined as a state and the path was the solution.



Hill-Climbing Search aka Greedy Local Search

Idea: keep a single “current” state and try to find better neighboring states.

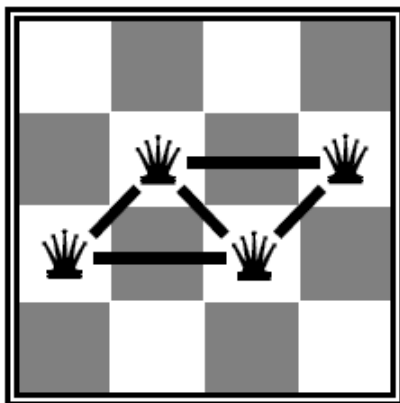
Example: n -queens problem

- **Goal:** Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal.
- **State space:** all possible n -queen configurations. We can restrict the state space: Only one queen per column.
- **Objective function:** minimize the number of pairwise conflicts.

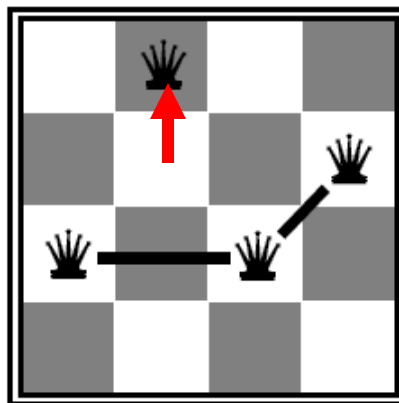
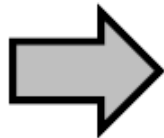
State space is reduced from 1820 to $4^4 = 256$

What is a possible local improvement strategy?


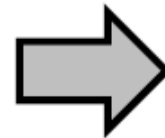
- Move one queen within its column to reduce conflicts



$h = 5$



$h = 2$



$h = 0$

Example: n -queens problem

- **Goal:** Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal.
- **State space:** all possible n -queen configurations. We can restrict the state space: Only one queen per column.
- **Objective function:** minimize the number of pairwise conflicts.

What is a possible local improvement strategy?

- Move one queen within its column to reduce conflicts

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	13	13	16	13	16
17	14	17	15	15	14	16	16
17	14	16	18	15	15	15	16
18	14	15	15	14	14	16	16
14	14	13	17	12	14	12	18

$h = 17$

best local improvement has $h = 12$

Note that there are many options, and we have to choose one!


Example: n -queens problem

Optimization problem: find the best arrangement a



$$a^* = \operatorname{argmin}_a \operatorname{conflicts}(a)$$

s.t. a has one queen per column



Remember: This makes the problem a lot easier.



$h = 5$



$h = 2$



$h = 0$

Example: Traveling Salesman Problem


- **Goal:** Find the shortest tour connecting n cities
- **State space:** all possible tours
- **Objective function:** length of tour




What's a possible local improvement strategy?

- Start with any complete tour, perform pairwise exchanges.

Permutation: **ABDEC**



ABCED



Example: Traveling Salesman Problem


Optimization problem: Find the best tour π

$$\pi^* = \operatorname{argmin}_{\pi} \text{tourLength}(\pi)$$


s.t. π is a valid permutation (i.e., sub-tour elimination)



Permutation: **ABDEC**



ABCED



Hill-climbing search (= Greedy local search)

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  current  $\leftarrow$  problem.INITIAL
  while true do
    neighbor  $\leftarrow$  a highest-valued successor state of current
    if VALUE(neighbor)  $\leq$  VALUE(current) then return current
    current  $\leftarrow$  neighbor
```

Variants:

- **Steepest-ascend hill climbing**
 - Check all possible successors and choose the highest-valued successors.
- **Stochastic hill climbing**
 - choose randomly among all uphill moves, or
 - generate randomly one new successor at a time until a better one is found = first-choice hill climbing – the most popular variant, this is what people often mean when they say “stochastic hill climbing”


Random-restart hill climbing – to deal with local optima

Hill-climbing search

Hill-climbing search is similar to greedy best-first search and the objective function as a (maybe not admissible) heuristic.

Is it complete/optimal?


- No – can get stuck in local optima



Example: local optimum for the 8-queens problem. No single queen can be moved within its column to improve the objective function.

$$h = 1$$

The state space “landscape”



How to escape local maxima?

→ Random restart hill-climbing can help.

What about “shoulders” (called “ridges” in higher dimensional space)?

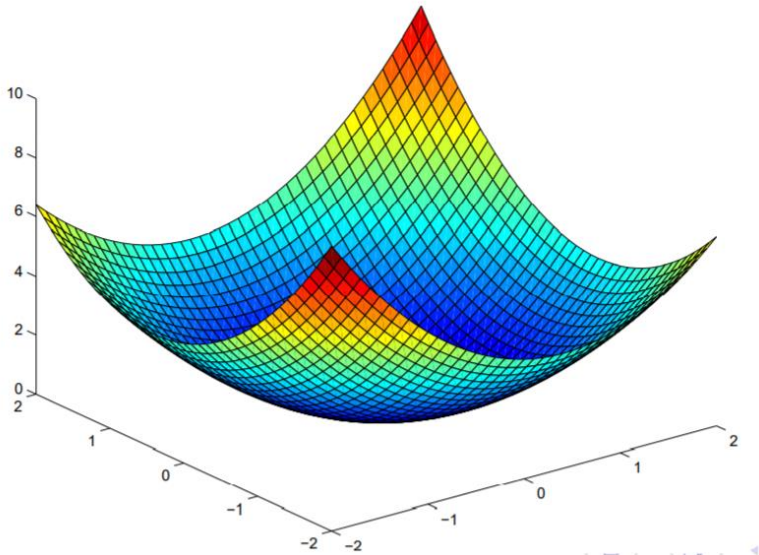
What about “plateaus”?

→ Hill-climbing with sideways moves.

Convex vs. Non-Convex Optimization Problems

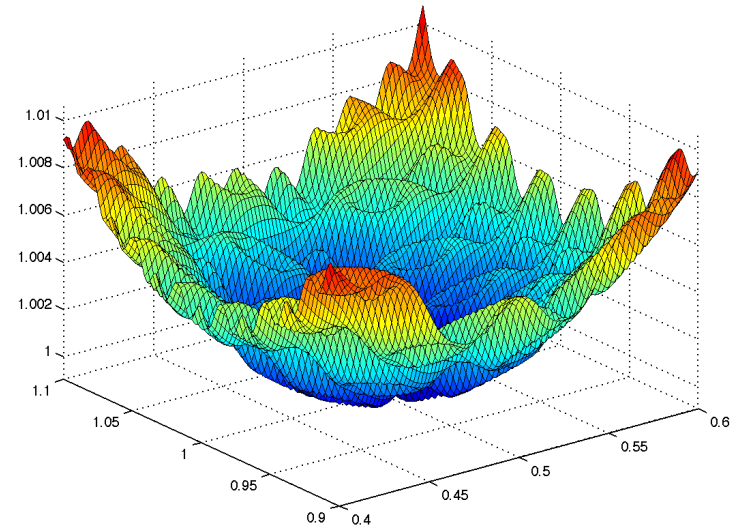
Minimization problem

Convex Problem



One global optimum +
smooth function → calculus
makes it easy

Non-convex Problem



Many local optima → hard

Many discrete optimization
problems are like this.


Minimization vs. Maximization

- The name **hill climbing** implies **maximizing a function**.
- Optimizers like to state problems as **minimization problems** (e.g., gradient descent).
- Both types of problems are equivalent:

$$\max(f(x))$$



$$\min(-f(x))$$



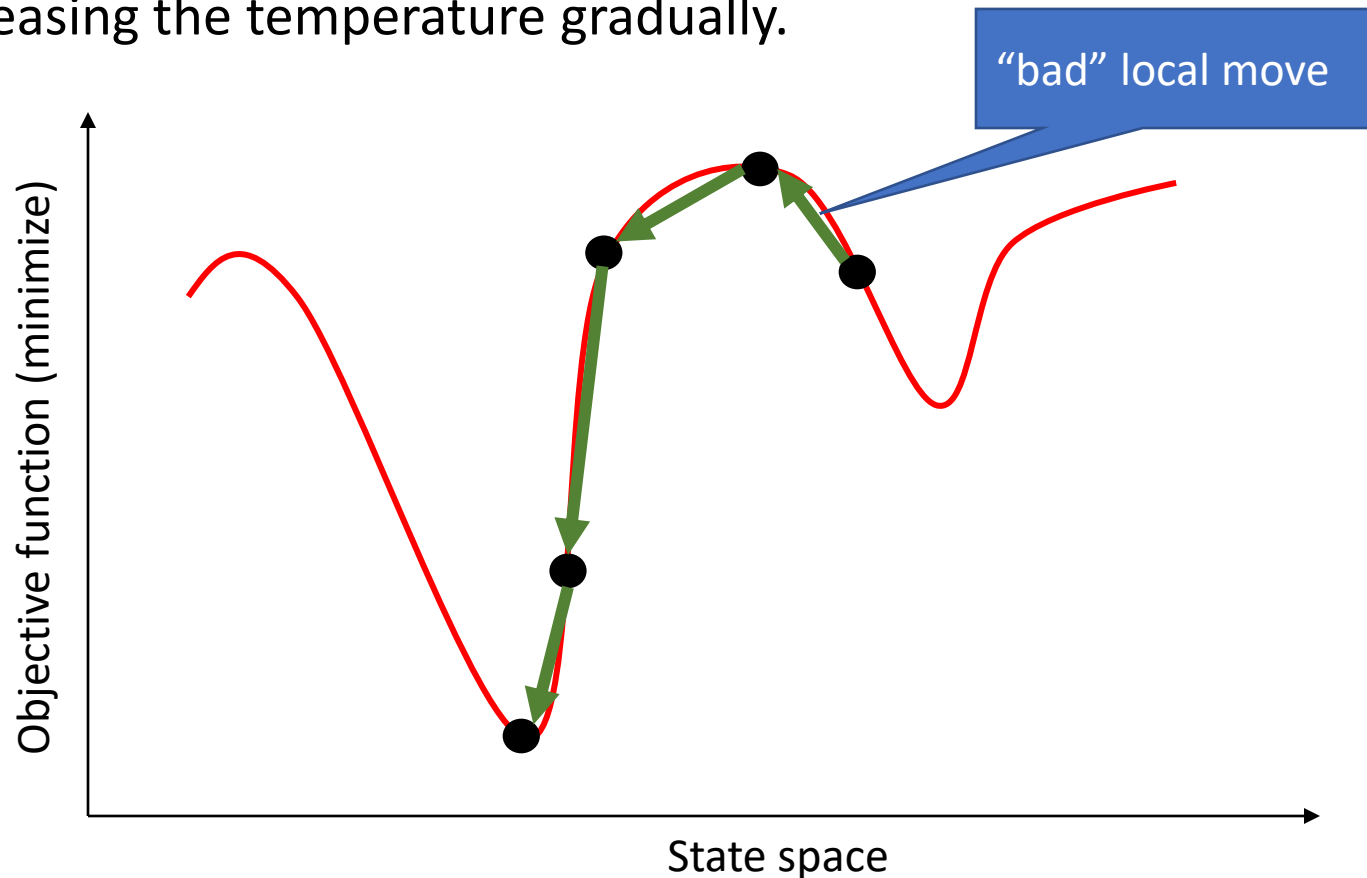
A glowing yellow-orange metal bar is being processed in a dark industrial machine. The bar is long and rectangular, with a bright, intense glow along its length, suggesting it is being heated or melted. The machine components are dark and metallic, with various rollers and guides visible. The background is dark, emphasizing the bright light of the metal.

Simulated Annealing

Use heat to escape local optima...

Simulated Annealing

- **Idea:** First-choice stochastic hill climbing + escape local minima by **allowing some “bad” moves** but gradually decrease their frequency.
- Inspired by the process of tempering or hardening metals by decreasing the temperature gradually.



Simulated Annealing

- **Idea:** First-choice stochastic hill climbing + escape local minima by allowing some “bad” moves but gradually decreasing their frequency as we get closer to the solution.
- The probability of accepting “bad” moves follows **an annealing schedule** that reduces the temperature T over time t .


```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state  
  current  $\leftarrow$  problem.INITIAL  
  for  $t = 1$  to  $\infty$  do  
     $T \leftarrow$  schedule( $t$ )  
    if  $T = 0$  then return current  
    next  $\leftarrow$  a randomly selected successor of current  
     $\Delta E \leftarrow$  VALUE(next) – Value(current)  
    if  $\Delta E > 0$  then current  $\leftarrow$  next  
    else current  $\leftarrow$  next only with probability  $e^{-\Delta E/T}$ 
```

Always do good moves

Uses the Metropolis acceptance criterion to accept “bad” moves

Note: Use VALUE(current) – VALUE(next) for minimization

The Effect of Temperature



The lower the temperature, the less likely the algorithm will accept a worse state.

Cooling Schedule

The cooling schedule is very important.

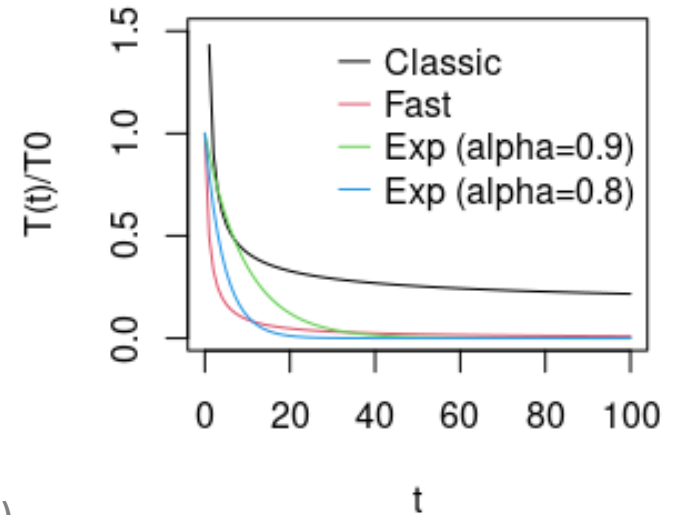
Popular schedules for the temperature at time t :

- **Classic simulated annealing:** $T_t = T_0 \frac{1}{\log(1+t)}$
- **Fast simulated annealing** (Szy and Hartley; 1987)

$$T_t = T_0 \frac{1}{1+t}$$

- **Exponential cooling** (Kirkpatrick, Gelatt and Vecchi; 1983)

$$T_t = T_0 \alpha^t \quad \text{for } 0.8 < \alpha < 1$$

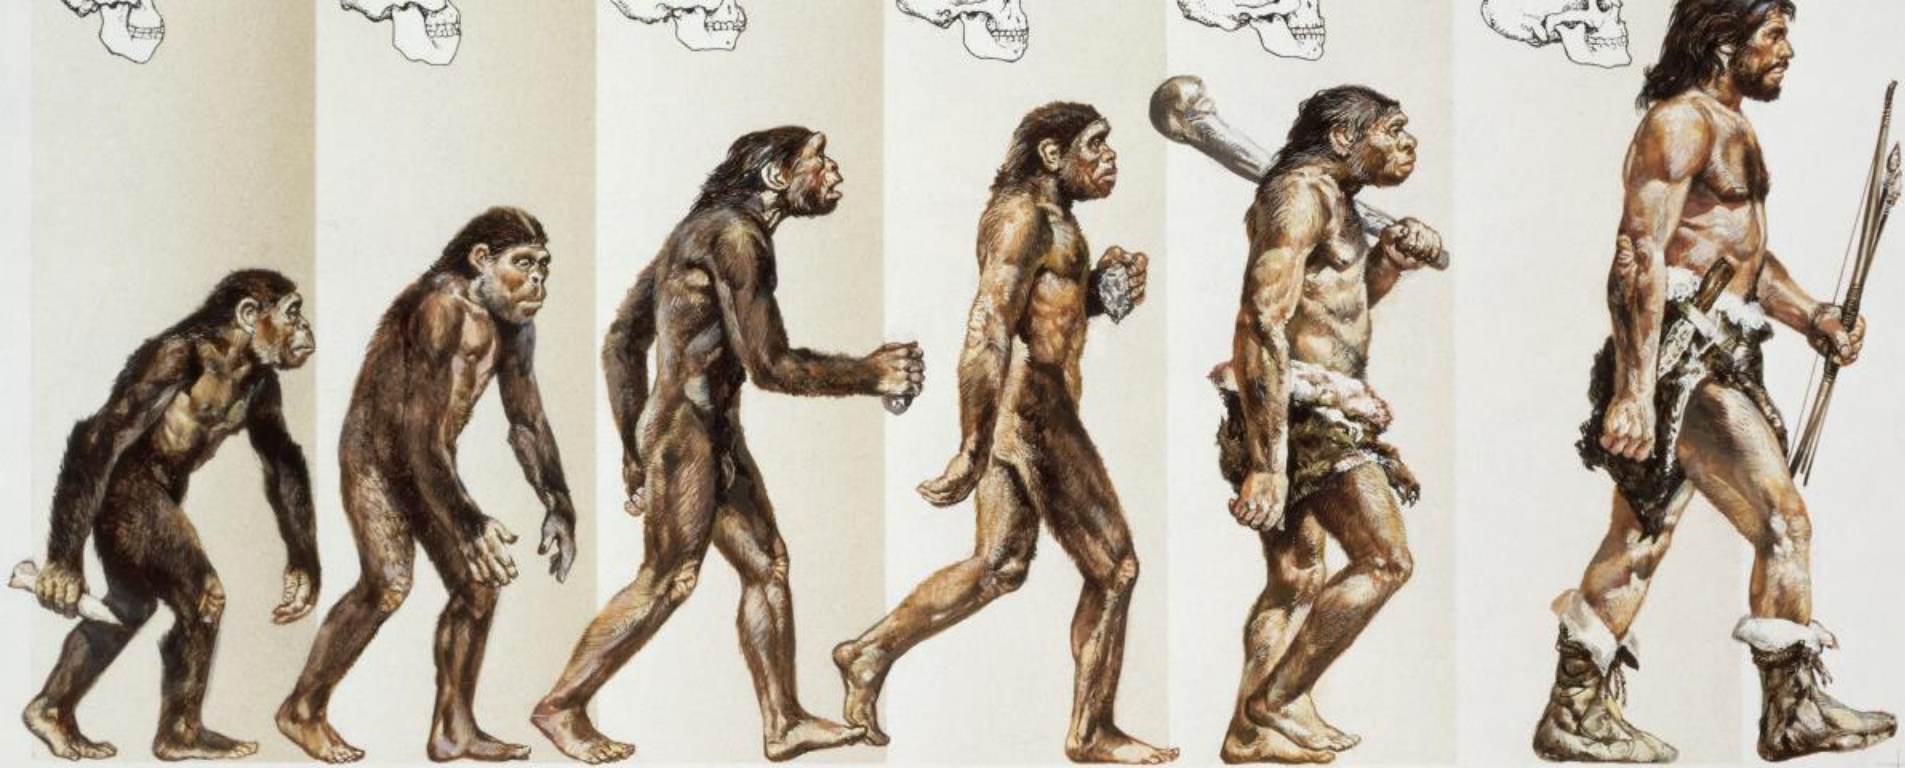


Notes:

- The best schedule is typically determined by trial-and-error.
- Choose T_0 to provide a high probability that any move will be accepted at time $t = 0$.
- T_t will not become 0 but very small. Stop when $T < \epsilon$ (ϵ is a very small constant).

Simulated Annealing Search

- **Guarantee:** If temperature decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching one.
- However:
 - This usually takes impractically long.
 - The more downhill/uphill steps you need to escape a local optimum, the less likely you are to make all of them in a row.
- The related **Markov Chain Monte Carlo (MCMC)** method is a general family of randomized algorithms for exploring complicated state spaces.



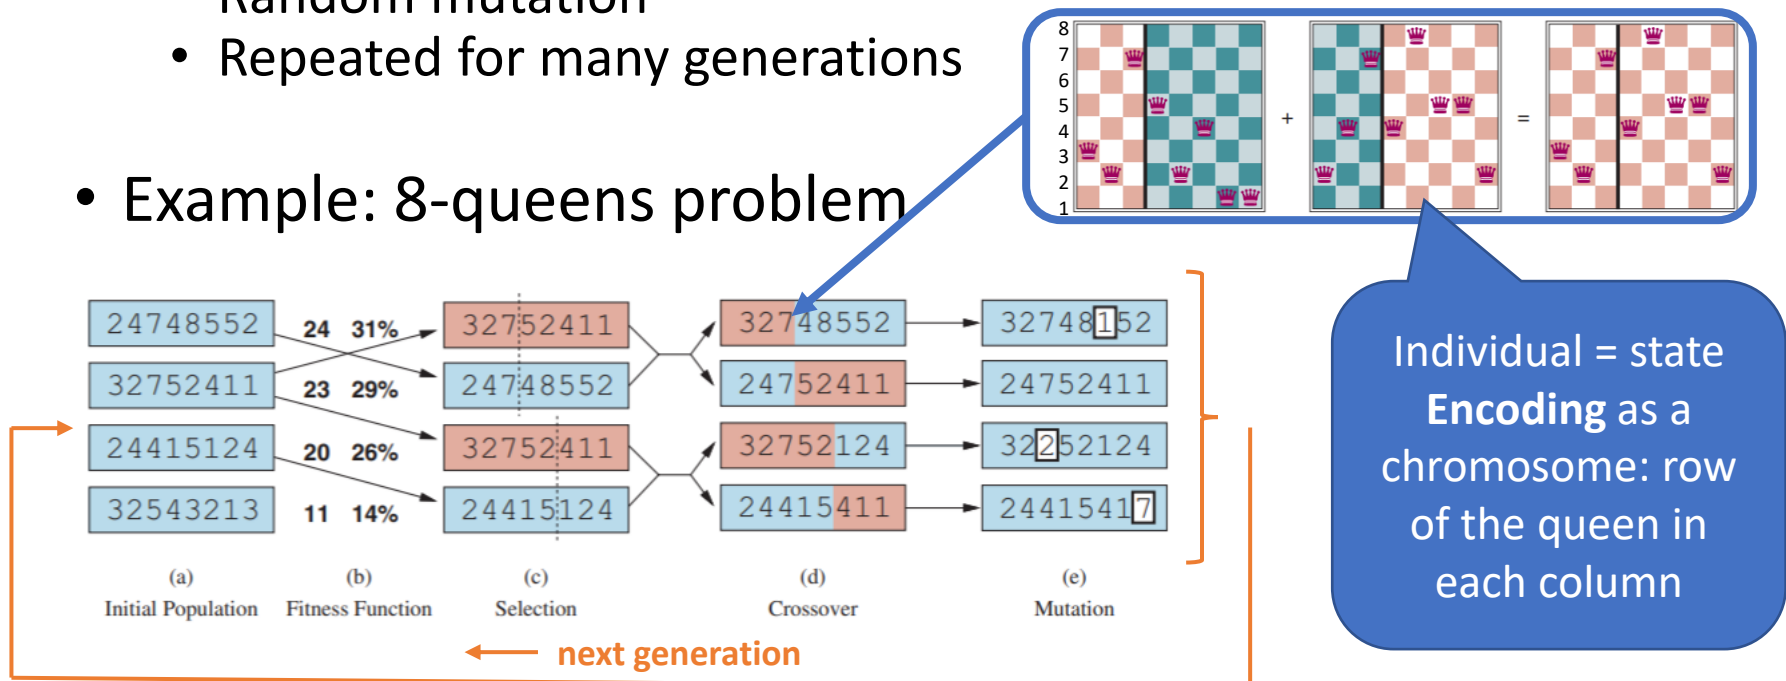
Evolutionary Algorithms


A Population-based Metaheuristics

Evolutionary Algorithms / Genetic Algorithms

- A metaheuristic for **population**-based optimization.
- Uses mechanisms inspired by biological evolution (genetics):
 - Reproduction: Random selection with probability based on a **fitness** function.
 - Random recombination (crossover)
 - Random mutation
 - Repeated for many generations

- Example: 8-queens problem

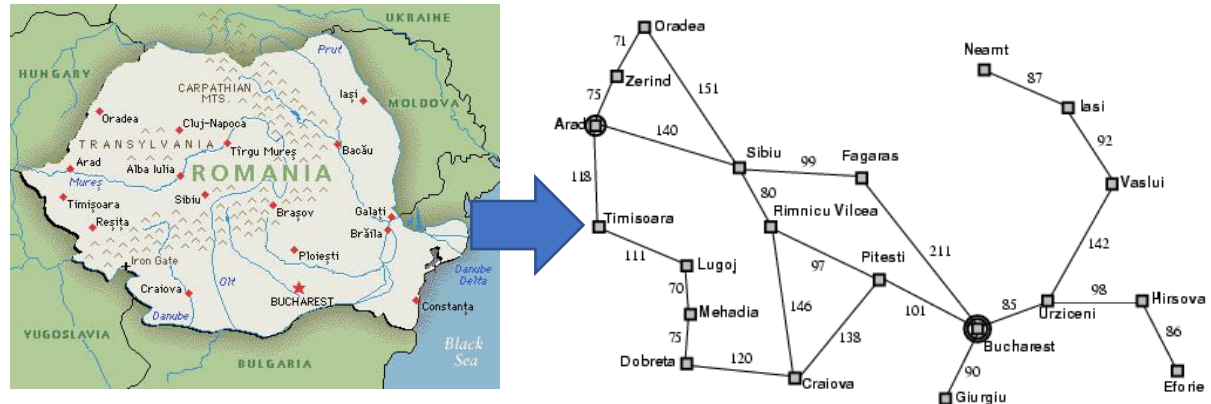





Search in Continuous Spaces

Discretization of Continuous Space

- Use atomic states and create a graph as the transition function.




- Use a grid with spacing of size δ
Note: You probably need a way finer grid!



Discretization of Continuous Space

How did we discretize this space?

Initial state



Goal
state

Search in Continuous Spaces: Gradient

Minimize $f(\mathbf{x}) = f(x_1, x_2, \dots, x_k)$

Gradient at point \mathbf{x} : $\nabla f(\mathbf{x}) = \left(\frac{\partial f(\mathbf{x})}{\partial x_1}, \frac{\partial f(\mathbf{x})}{\partial x_2}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_k} \right)$
(=evaluation of the Jacobian matrix at \mathbf{x})

Find optimum by solving: $\nabla f(\mathbf{x}) = 0$

- **Gradient descent (= Steepest-ascent hill climbing for minimization)**
with step size α



$$\mathbf{x} \leftarrow \mathbf{x} - \alpha \nabla f(\mathbf{x})$$

- **Newton-Raphson method**

uses the inverse of the Hessian matrix of the second derivative $H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$ for the step size α

$$\mathbf{x} \leftarrow \mathbf{x} - \mathbf{H}_f^{-1}(\mathbf{x}) \nabla f(\mathbf{x})$$

Note: May get stuck in a local optimum if the search space is non-convex! Use simulated annealing.



Search in Continuous Spaces: Empirical Gradient Methods

- What if the mathematical formulation of the objective function is not known?
 - We may have objective values at fixed points, called the training data.
 - In this case, we can use **empirical gradient search**. This is related to steepest ascend hill climbing in the discretized state space.
- We will talk more about search in continuous spaces with loss functions using gradient descend when we talk about **parameter learning for machine learning**.